


Safe, Fast, and Easy

Building a plugin system with WebAssembly

Kyle Conroy | conroy.org | @kyle_conroy



sqlc

sqlc

github.com/kyleconroy/sqlc | sqlc.dev

Compile SQL to type-safe code

How it works

How it works

1. You write SQL queries

How it works

1. You write SQL queries
2. You run `sqlc generate`, which outputs Go code with type-safe interfaces to those queries

How it works

1. You write SQL queries
2. You run `sqlc generate`, which outputs Go code with type-safe interfaces to those queries
3. You write application code that calls these methods


Quick example

```
CREATE TABLE authors (  
    id    BIGSERIAL PRIMARY KEY,  
    name  text          NOT NULL,  
    bio   text  
);
```

Quick example

```
-- name: GetAuthor :one  
SELECT * FROM authors  
WHERE id = $1 LIMIT 1;
```

Quick example



Quick example

```
ctx := context.Background()
queries := example.New(db) // *db.Sql

author, _ := queries.GetAuthor(ctx, 42)
fmt.Println(author.Name)
```

Column expansion

```
-- name: GetAuthor :one  
SELECT * FROM authors  
WHERE id = $1 LIMIT 1;
```

Column expansion

```
const getAuthQuery = `
SELECT id, name, bio FROM authors
WHERE id = $1 LIMIT 1;
`
```

Type inference

```
-- name: GetAuthor :one  
SELECT * FROM authors  
WHERE id = $1 LIMIT 1;
```


Type inference

```
func (q *Queries) GetAuthor(  
    ctx context.Context,  
    id int64,  
) (Author, error) {  
    ...  
}
```

Catches typos

```
-- name: GetAuthor :one  
SELECT first_name FROM authors  
WHERE id = $1 LIMIT 1;
```

Advanced SQL

- Common Table Expressions
- Extensions
- Type inference for built-in functions
- Enums
- DDL

Supported databases

- PostgreSQL
- MySQL
- SQLite (beta)

GOPHERCON

2022



SQL HERO CON 2022



SQL HERO CON 2022






Compile SQL to type-safe code

Creating magic with ASTs and parsers

Kyle Conroy | conroy.org | @kyle_conroy



Extending sqlc

Extending sqlc

- Changing existing codegen output
- Adding database engines
- Adding programming languages

Extending sqlc


- **Changing existing codegen output**
 - Configuration file grew from a few options to over thirty
- Adding database engines
- Adding programming languages

Extending sqlc

- Changing existing codegen output
- **Adding database engines**
 - Started with PostgreSQL but added MySQL a bit afterwards
- Adding programming languages

Extending sqlc

- Changing existing codegen output
- Adding database engines
- **Adding programming languages**
 - Early requests to merge support for Kotlin and Python






```
graph LR; A[parser] --> B[compiler]; B --> C[codegen];
```


parser

compiler

codegen




Starting point



`codegen.Generate()`




`map[string]string{}`




Starting point

- The codegen package only knew how to generate Go
- Imported a ton of internal shared packages
- Mainly implemented via text/template




Go packages



```
interface Generator {  
    Generate()  
}
```



```
map[string]string{}
```




Go packages

- fir

Go packages

- Technically worked, merged Kotlin and Python support
- Many downsides
- Maintenance headaches



Requirements

Requirements

Context

- Command line tool, doesn't control the host environment
- Needs to support Linux, macOS, and Windows
- Plugins for specific languages are best written in that language

Requirements

- Independent
- Safe
- Run anywhere
- Fast
- Familiar

Requirements

- Independent
- Safe
- Run anywhere
- Fast
- Familiar

Requirements

- Independent
- Safe
- Run anywhere
- Fast
- Familiar

Requirements

- Independent
- Safe
- Run anywhere
- Fast
- Familiar

Requirements

- Independent
- Safe
- Run anywhere
- Fast
- Familiar

Requirements

- Independent
- Safe
- Run anywhere
- Fast
- Familiar

Go packages

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Go packages

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Go packages

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Go packages


Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Go packages

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Attempts

- ~~Go packages~~



plugin

plugin

package

standard library

Version: [go1.19.2](#) **Latest** | Published: Oct 4, 2022 |

License: [BSD-3-Clause](#) | Imports: [4](#) | Imported by: [2,173](#)

plugin

package

standard library

Version: [go1.19.2](#) **Latest** | Published: Oct 4, 2022 |

License: [BSD-3-Clause](#) | Imports: [4](#) | Imported by: [2,173](#)


```
p := plugin.Open("codegen.so")  
f := p.Lookup("Generate")  
f()
```

compiler



```
map[string]string{}
```

codegen
(plugin)

plugin

- "Currently plugins are only supported on Linux, FreeBSD, and macOS"
 - No Windows!

plugin

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

plugin

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

evil_👿_plugin




- Tries to find AWS credentials on your system
 - Looks in ~/.aws/credentials
 - Looks for AWS_SECRET_KEY
- And then send them off to the Cloud™

evil_👿_plugin

"PyPi python packages caught sending stolen AWS keys to unsecured sites"

<https://www.bleepingcomputer.com/news/security/pypi-python-packages-caught-sending-stolen-aws-keys-to-unsecured-sites/>



plugin

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

plugin


Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

plugin


Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Attempts

- ~~Go packages~~
- plugin




os/exec



message CodeGenRequest



message CodeGenResponse



os/exec

- protoc uses this model
 - Lookup a plugin on \$PATH
 - Start a process, send data via STDIN and read from STDOUT
- Well-understood pattern and supported across many languages




os/exec

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				



os/exec

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				



os/exec

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

os/exec


Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

os/exec

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Attempts


- ~~Go packages~~
- plugin
- ~~os/exec~~




hashicorp/go-plugin




HashiCorp
Nomad



HashiCorp
Boundary




HashiCorp
Terraform



HashiCorp
Vault



HashiCorp
Packer




message CodeGenRequest



gRPC



message CodeGenResponse



hashicorp/go-plugin

- Plugins required to implement a large, complicated interface
 - At the end of the day it's just gRPC
 - I'm still not sure how easy it is to implement in a different language

hashicorp/go-plugin

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				




hashicorp/go-plugin

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

hashicorp/go-plugin

- "Plugins can be **relatively** secure: The plugin only has access to the interfaces and args given to it, not to the entire memory space of the process. Additionally, go-plugin can communicate with the plugin over TLS."
- Don't be fooled!


hashicorp/go-plugin

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

hashicorp/go-plugin

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

hashicorp/go-plugin


Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

hashicorp/go-plugin


- Obviously not a bad choice!
- Powers an ecosystem many, many times larger than sqlc
- Much easier to secure if you're running the plugins in the cloud

Attempts

- ~~Go packages~~
- plugin
- ~~os/exec~~
- ~~hashicorp/go-plugin~~



Embedded scripting




message CodeGenRequest



interpreter



message CodeGenResponse



Embedded scripting

- Lua / JavaScript / Python
- Have to find an interpreter that meets your needs

Embedded scripting

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

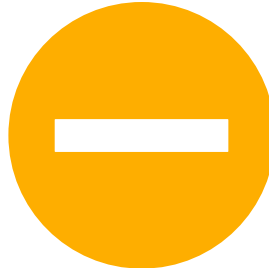

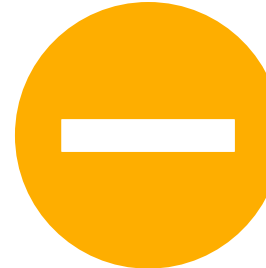
Embedded scripting

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

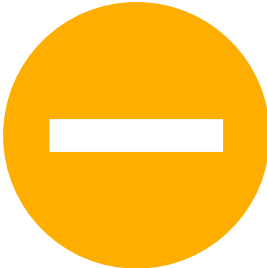
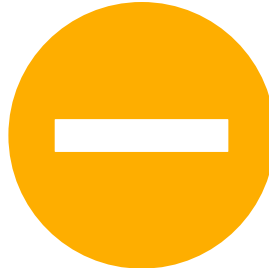



Embedded scripting

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Embedded scripting


Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Embedded scripting

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Attempts

- ~~Go packages~~
- plugin
- ~~os/exec~~
- ~~hashicorp/go-plugin~~
- scripting



bytecodealliance/wasmtime-go

"A fast and secure runtime for
WebAssembly"


<https://wasmtime.dev/>

"WebAssembly (abbreviated *Wasm*) is a binary instruction format for a stack-based virtual machine. Wasm is designed as a portable compilation target for programming languages, enabling deployment on the web for client and server applications."

<https://webassembly.org/>

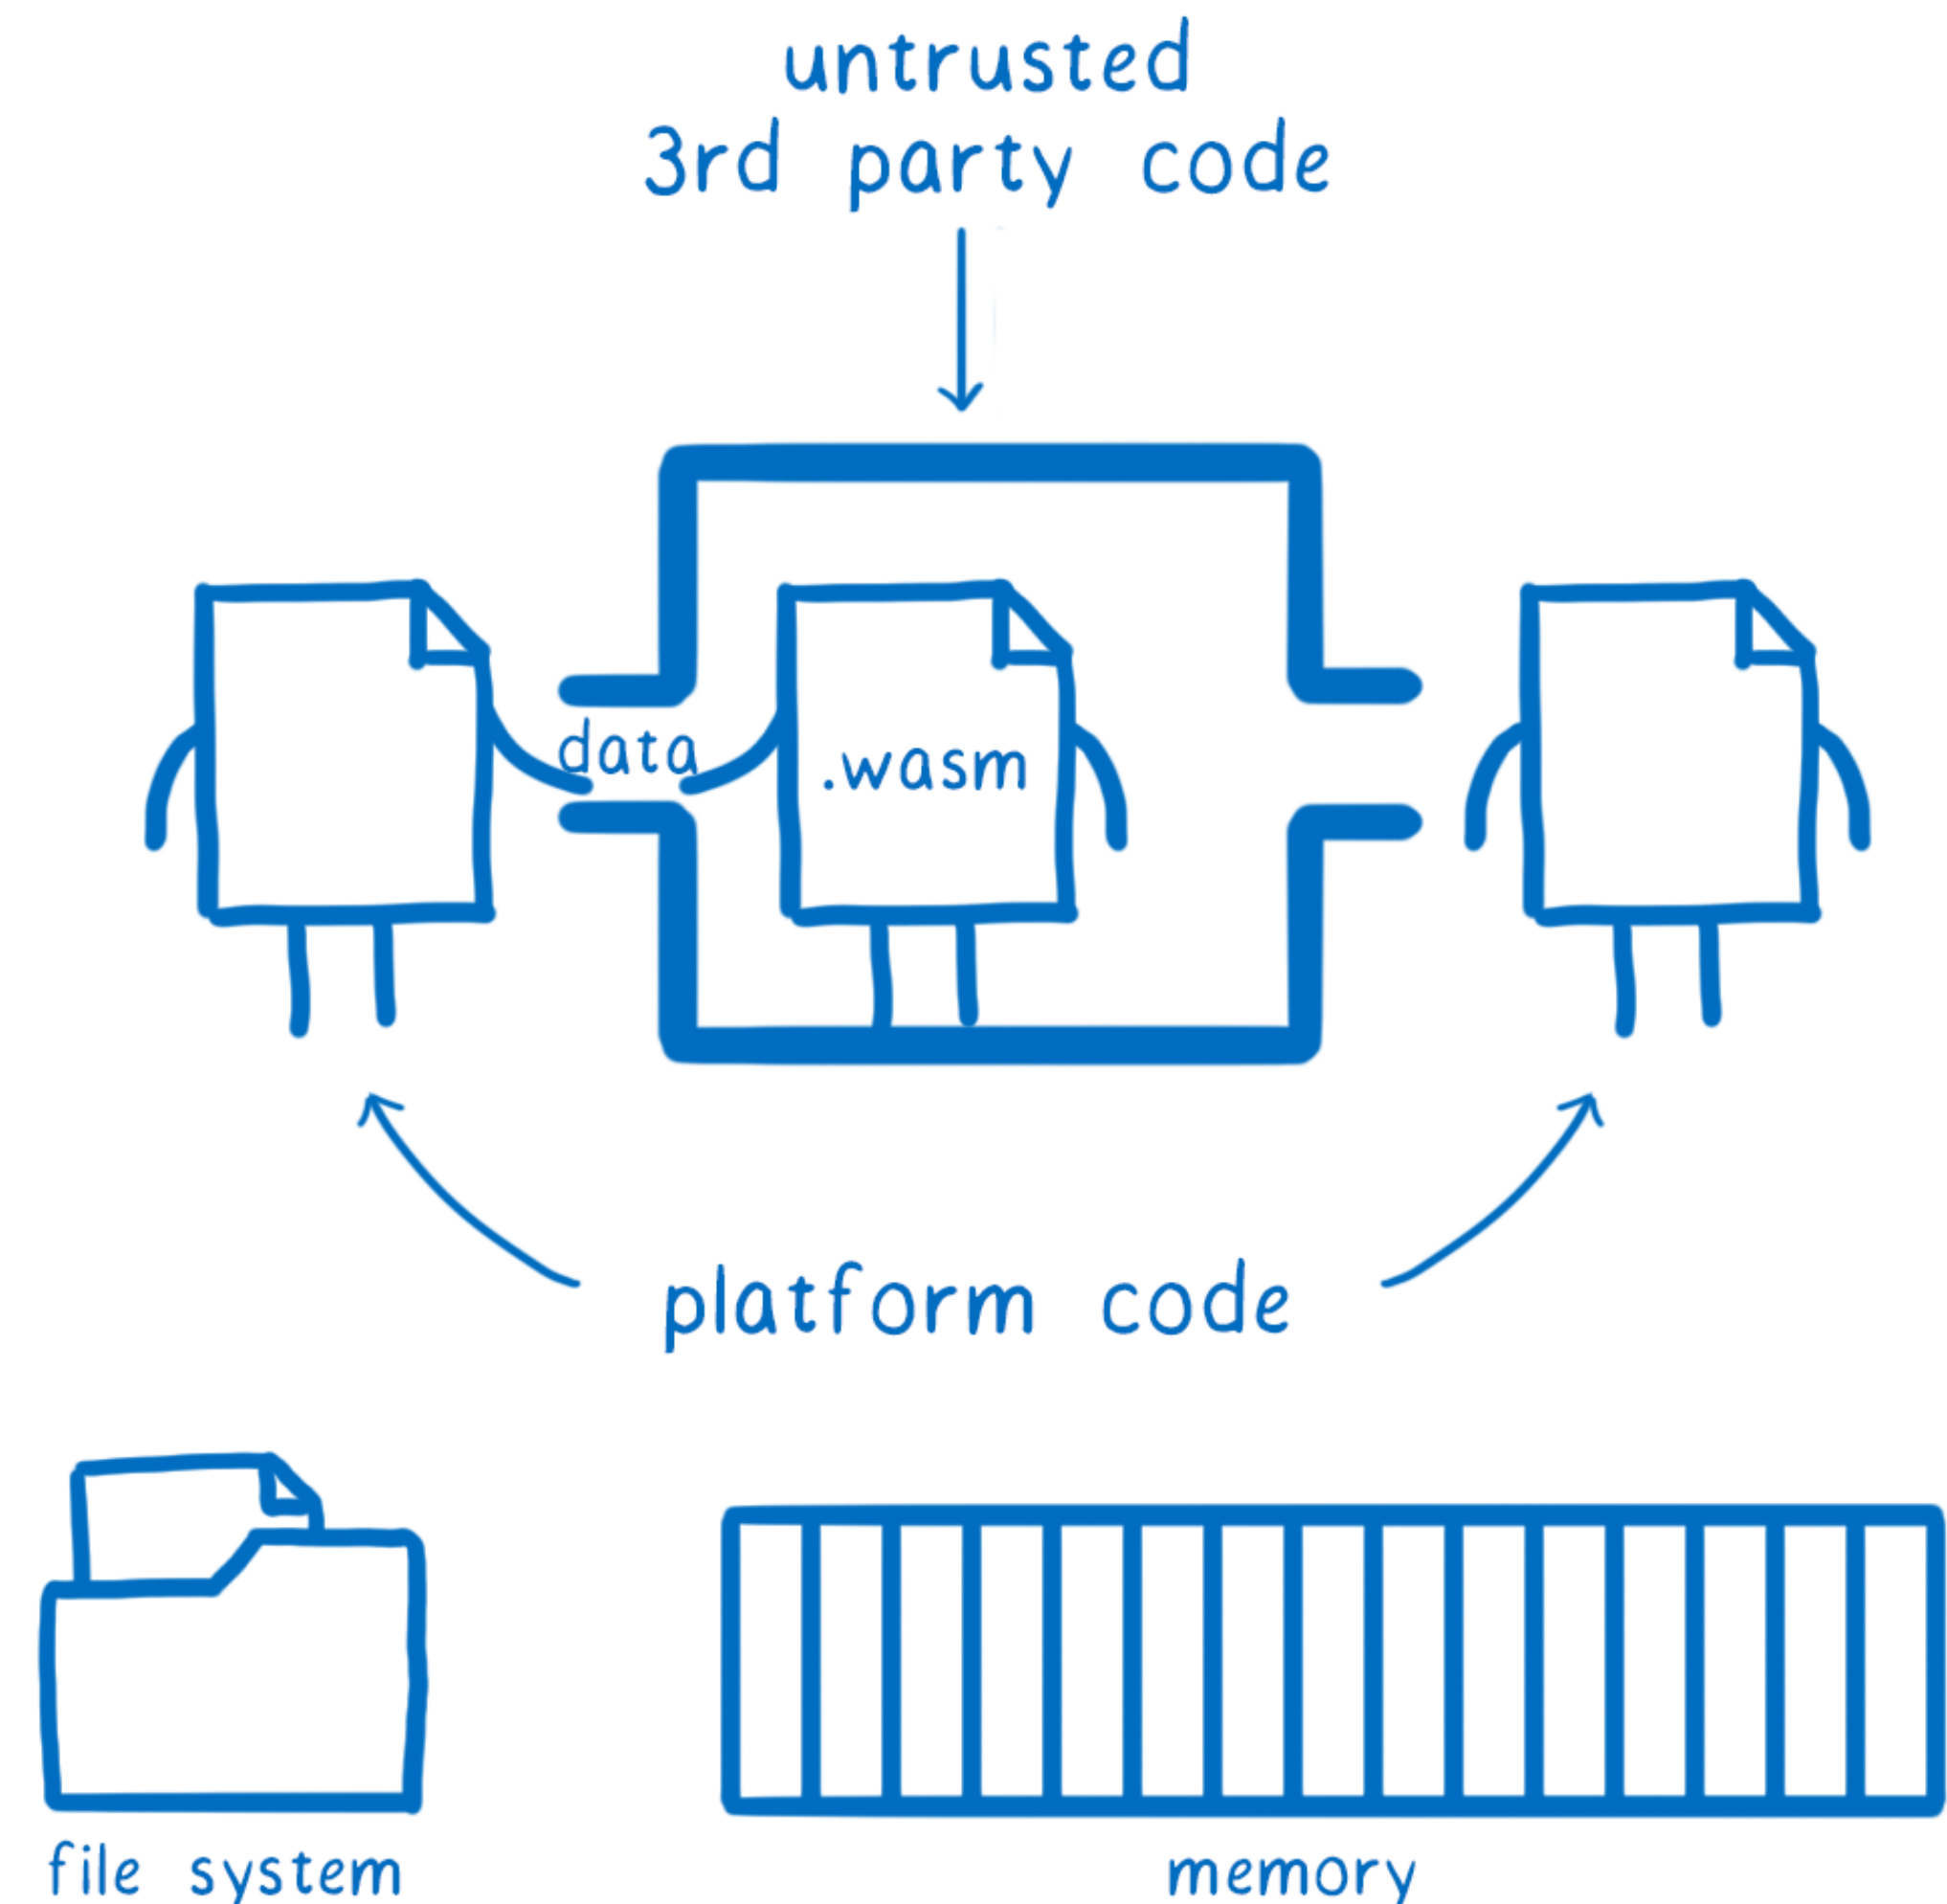
3rd Party Plugin systems

WebAssembly is great for platforms, where you often want to run 3rd party code so you can support many different specific use cases—for example, **plug-in** marketplaces where developers in the platform's **ecosystem** can share code with users."



3rd Party Plugin systems

WebAssembly is great for platforms, where you often want to run 3rd party code so you can support many different specific use cases—for example, **plug-in** marketplaces where developers in the platform's **ecosystem** can share code with users."



bytecodealliance/wasmtime-go

- Recently-released stable 1.0.0!
- Platform-independent modules!
- Secure by default!

bytecodealliance/wasmtime-go

Security


- Capability-based security system
- By default:
 - No filesystem access
 - No network access
 - No environment variables
- Can give access to specific directories

bytecodealliance/wasmtime-go

- Wasm is low-level; doesn't know about things like strings, objects, etc.
- Wasm-only interfaces are... intimidating

"WASI stands for WebAssembly System Interface. It's an API ... that provides access to several operating-system-like features, including files and filesystems..."

<https://wasi.dev/>



message CodeGenRequest




stdin



stdout

message CodeGenResponse




bytecodealliance/wasmtime-go

- We've actually gone full circle back to `exec/cmd`
- Wasm + WASI is the same model as process-based plugins

bytecodealliance/wasmtime-go

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				


bytecodealliance/wasmtime-go

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				




bytecodealliance/wasmtime-go

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

bytecodealliance/wasmtime-go


Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

bytecodealliance/wasmtime-go

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

Attempts

- ~~Go packages~~
- plugin
- ~~os/exec~~
- ~~hashicorp/go-plugin~~
- ~~scripting~~
- **wasmtime-go**



Implementation

Configuration

version: '2'

plugins:

- name: greeter

 wasm:

 url: <https://github.com/kyleconroy/.../sqlc-gen-greeter.wasm>

 sha256: "afc486dac206.....cd802424ad07"

Configuration

sql:

- schema: schema.sql

queries: query.sql

engine: postgresql

codegen:

- out: gen

plugin: greeter

How it works

1. Load the configuration file
2. Download each plugin's .wasm file (if necessary)
3. Create a wasmtime execution context (engine / module / linker)
4. Set up stdin, stdout, stderr
5. Call the context, read stdout

Speed

baseline	wasmtime-go
0.035s	0.955s

27x

slower

Speed (bump)



How it works (faster)




1. Load the configuration file
2. Download each plugin's .wasm file (if necessary)
- 3. Turn the .wasm file into a module, save to disk**
4. Create a wasmtime execution context (engine / linker)
5. Set up stdin, stdout, stderr
6. Call the context, read stdout

Speed (up)

baseline	wasmtime-go
0.035s	0.055s

`internal/ext/wasm/wasm.go`

bytecodealliance/wasmtime-go

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

bytecodealliance/wasmtime-go

Independent?	Secure?	Run anywhere?	Fast?	Familiar?
				

GOOS=js GOARCH=wasm

- Go's builtin Wasm support does not support WASI
- So off to TinyGo we... go!

TinyGo, big problems

- TinyGo doesn't work with common serialization tools
 - No encoding/json
 - No encoding/xml
 - No protobuf

GOPHERCON

2022







Crustaceans and Codegen

Building sqlc plugins in Rust

Kyle Conroy | conroy.org | @kyle_conroy

sqlc-gen-greeter

<https://github.com/kyleconroy/sqlc-gen-greeter>

- Written in Rust
- Extremely basic, just outputs "Hello world"

sqlc-gen-node-pg

<https://github.com/tabbed/sqlc-gen-node-pg>

- TypeScript output for sqlc
- Written in Rust
- Still a work-in-progress

Attempts

- ~~Go packages~~
- plugin
- ~~os/exec~~
- ~~hashicorp/go-plugin~~
- ~~scripting~~
- **wasmtime-go**

Attempts

- ~~Go packages~~
- plugin
- **os/exec**
- ~~hashicorp/go-plugin~~
- ~~scripting~~
- **wasmtime-go**

sqlc-gen-json

<https://github.com/kyleconroy/sqlc/tree/main/cmd/sqlc-gen-json>

- Process-based plugin written in Go
- Not all languages support Wasm and WASI
- No distribution built-in (on purpose)

2

asks

Serialize and deserialize protocol buffers

<https://github.com/tinygo-org/tinygo/issues/2667>

Go toolchain support for WASI

Thanks!

- Twitter: @kyle_conroy
- Email: kyle@conroy.org