

CGO-less Foreign Function Interface with WebAssembly

GopherCon 2022 - Takeshi Yoneda at Tetrade



Foreign Function Interface(FFI)



Foreign Function Interface(FFI)

*“A foreign function interface (FFI) is a mechanism by which a program written in **one programming language can call routines or make use of services written in another.**” – wikipedia*



```
func main () {
```

```
}
```



main.go



```
func main () {  
    rustFn()
```

```
}
```



main.go



```
pub extern "C" fn rustFn() {...}
```

lib.rs



```
func main () {  
    rustFn()  
    zigFn()  
    ....  
}
```



main.go

```
pub extern "C" fn rustFn() {...}
```

lib.rs



```
export fn zigFn() void { ... }
```



lib.zig



When/Why do we want FFI?



When/Why do we want FFI?

- Reusing softwares in other languages
 - Don't want to rewrite 100k loc in C



When/Why do we want FFI?

- Reusing softwares in other languages
 - Don't want to rewrite 100k loc in C
- Plugin System via FFI - Polyglot!
 - Allow users to extend your app in any language



```
func main () {  
    rustFn()  
    zigFn()  
    ....  
}
```



main.go

```
pub extern "C" fn rustFn() {...}
```



lib.rs

```
export fn zigFn() void { ... }
```




lib.zig



```
func main () {  
    b := rustFn(v)  
    b = zigFn(v)  
    ....  
}
```



main.go




What's the protocol between Go and another lang?

How Go runtime behaves beyond Go world?



What's the proto and another lang?

How Go runtime and Go world?



CGO



```
func main () {
    b := rustFn(v) {}
    b = zigFn(v) {}
    ....
}
```



main.go



CGO

```
pub extern "C" fn rustFn(v: u32) -> bool {...}
```



lib.rs

```
export fn zigFn(v: u32) bool { ... }
```



lib.zig



FFI can be done with CGO. The problem solved?

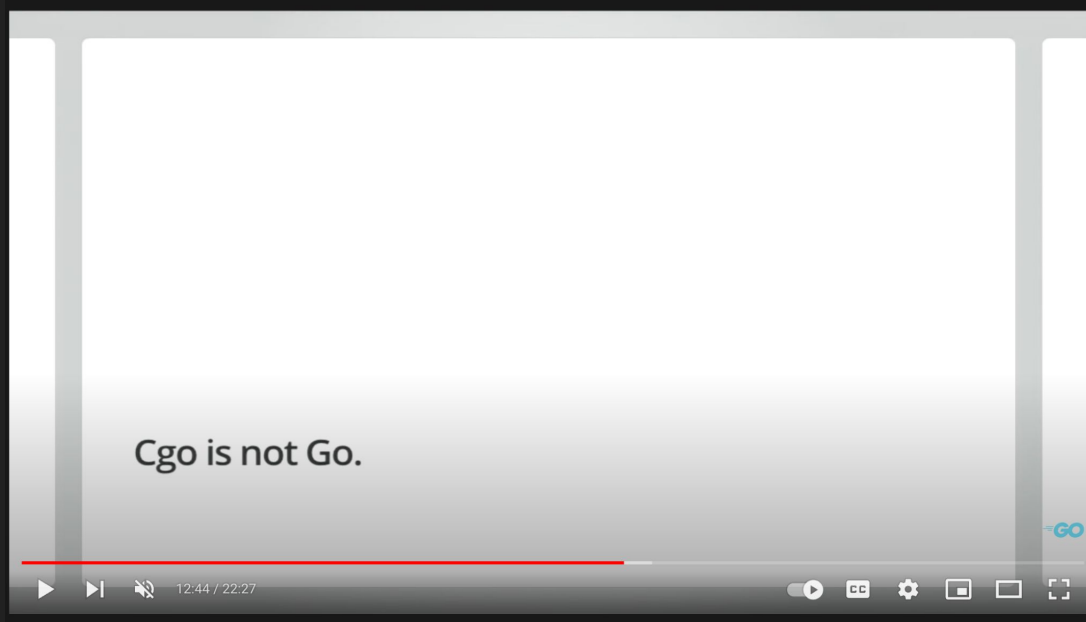


FFI can be done with CGO. The problem solved?

No.



“CGO is not Go”



Gopherfest 2015 | Go Proverbs with Rob Pike <https://youtu.be/PAAkCSZUG1c>



CGO troubles

- Dynamic vs Static binary: portability issue
- Cross compilation
- CGO is slow
- Security



CGO troubles

- Dynamic vs Static binary: portability issue
- Cross compilation
- CGO is slow
- Security



```
package main
```

```
func main() { println("hello") }
```



```
package main
```

```
func main() { println("hello") }
```

```
$ go build main.go
```



```
package main
```

```
func main() { println("hello") }
```

```
$ go build main.go
```

```
$ ldd main
```



```
package main
```

```
func main() { println("hello") }
```

```
$ go build main.go
```

```
$ ldd main
```

```
not a dynamic executable
```



```
package main
```

```
import "C"
```

```
func main() { println("hello") }
```




```
package main
```

```
import "C"
```

```
func main() { println("hello") }
```

```
$ go build main.go
```



```
package main
```

```
import "C"
```

```
func main() { println("hello") }
```

```
$ go build main.go
```

```
$ ldd main
```



```
package main
```

```
import "C"
```

```
func main() { println("hello") }
```

```
$ go build main.go
```

```
$ ldd main
```

```
linux-vdso.so.1 (0x00007ffd59db2000)
```

```
libpthread.so.0 => /lib/x86_64-linux-gnu/libpthread.so.0 (0x00007fad32c3f000)
```

```
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fad32a4d000)
```

```
/lib64/ld-linux-x86-64.so.2 (0x00007fad32c83000)
```



CGO troubles

- Dynamic vs Static binary
- Cross compilation
- CGO is slow
- Security



```
~/hugo master
>> uname
Darwin

~/hugo master
>> CGO_ENABLED=1 GOOS=darwin go build ./...

~/hugo master
>> CGO_ENABLED=1 GOOS=linux go build ./...
# runtime/cgo
linux_syscall.c:67:13: error: implicit declaration of function 'setresgid' is invalid in C99 [-Werror,-Wimplicit-function-declaration]
linux_syscall.c:67:13: note: did you mean 'setregid'?
/Applications/Xcode.app/Contents/Developer/Platforms/MacOSX.platform/Developer/SDKs/MacOSX12.3.sdk/usr/include/unistd.h:593:6: note: 'setregid' declared here
linux_syscall.c:73:13: error: implicit declaration of function 'setresuid' is invalid in C99 [-Werror,-Wimplicit-function-declaration]
linux_syscall.c:73:13: note: did you mean 'setreuid'?
/Applications/Xcode.app/Contents/Developer/Platforms/MacOSX.platform/Developer/SDKs/MacOSX12.3.sdk/usr/include/unistd.h:595:6: note: 'setreuid' declared here
```



CGO troubles

- Dynamic vs Static binary
- Cross compilation
- CGO is slow
- Security



runtime: performance problem with many Cgo calls #19574

New issue

Open

petermattis opened this issue on Mar 17, 2017 · 14 comments



petermattis commented on Mar 17, 2017



Please answer these questions before submitting your issue. Thanks!

What version of Go are you using (`go version`)?

`go version go1.8 darwin/amd64`, though I've also verified this happens on Linux.

What operating system and processor architecture are you using (`go env`)?

Assignees

No one assigned

Labels

compiler/runtime

Performance

Projects

None yet

<https://github.com/golang/go/issues/19574>



CGO troubles

- Dynamic vs Static binary
- Cross compilation
- CGO is slow
- Security




```
func main () {  
  b := rustFn(v)  $\longleftrightarrow$   
  b = zigFn(v)  $\longleftrightarrow$   
  ....  
}
```



CGO

pub extern "C" fn rustFn(v: u32) -> bool {...}



export fn zigFn(v: u32) bool {...}



Operating System




```
func main () {
  b := rustFn(v)
  b = zigFn(v)
  ....
}
```



CGO

```
pub extern "C" fn rustFn(v: u32) -> bool {...}
```



```
export fn zigFn(v: u32) bool {...}
```



Operating System



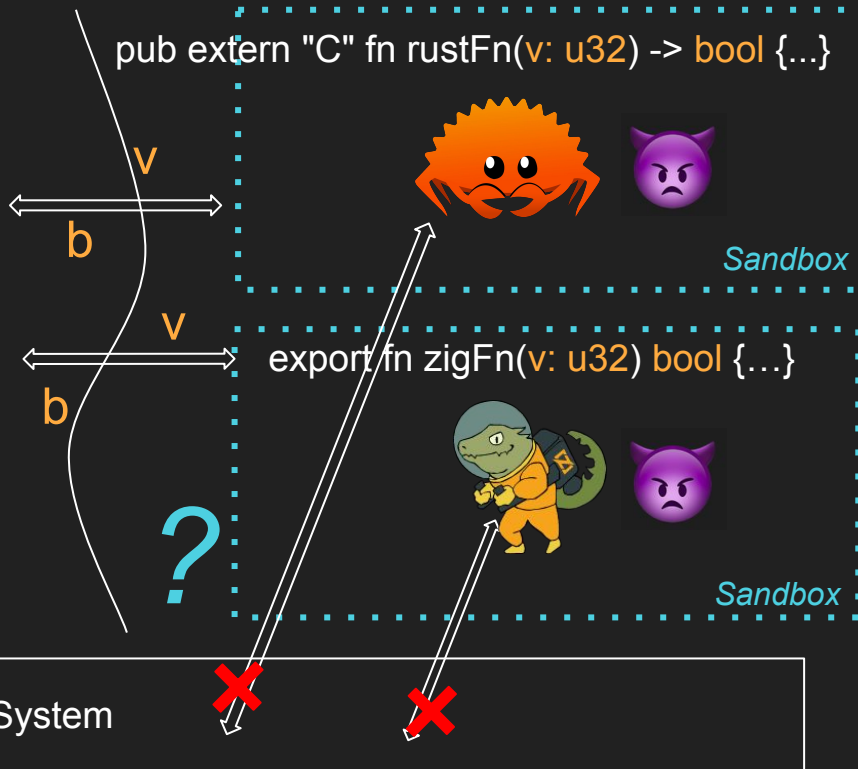
We need sandox...



```
func main () {  
  b := rustFn(v)  $\longleftrightarrow$   
  b = zigFn(v)  $\longleftrightarrow$   
  ....  
}
```



?





WebAssembly (Wasm)



WebAssembly (Wasm)

- Binary instruction format for a stack-based virtual machine (VM)
- Polyglot
- Security-oriented design
 - Memory guard
 - Deny system calls by default



Wait, isn't WebAssembly for the web?



Wasm is not only for the browsers

- Core spec is decoupled from the web concept
- Embeddable in any application with VM implementation



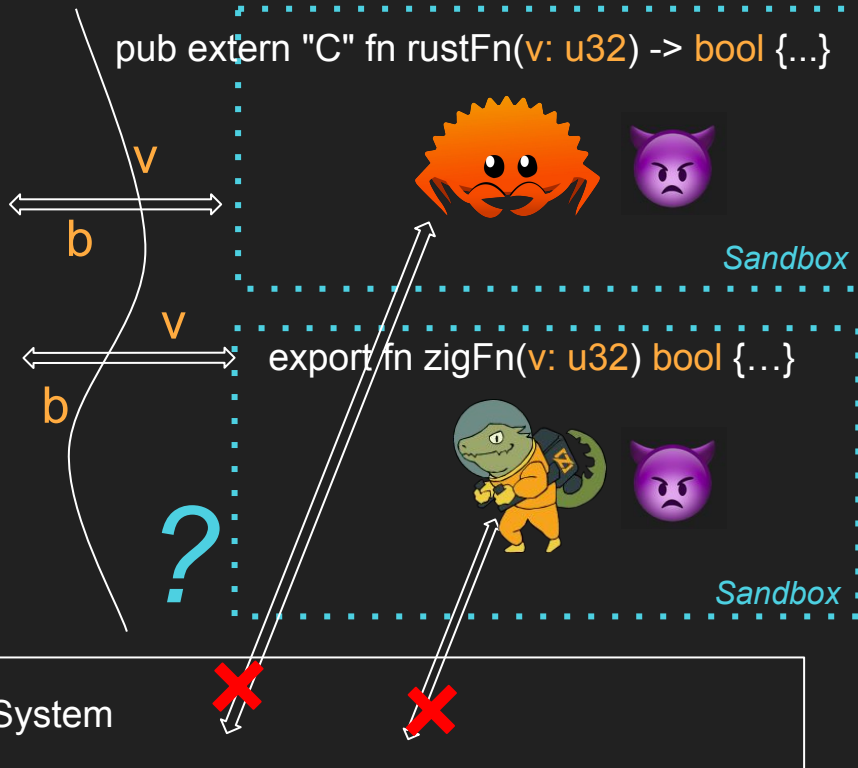
Non-Web examples




```
func main () {  
  b := rustFn(v)  $\longleftrightarrow$   
  b = zigFn(v)  $\longleftrightarrow$   
  ....  
}
```



?



```
func main () {  
  b := rustFn(v)  $\longleftrightarrow$   
  b = zigFn(v)  $\longleftrightarrow$   
  ....  
}
```



How to run Wasm binary inside Go?



Wasm needs a VM runtime!



x86_64
aarch64
riscv64
....





wazero.io

the zero dependency WebAssembly runtime for Go developers



What is wazero?


- Started out as my hobby project: now sponsored by Tetrate
- The Wasm runtime with zero dependency
- Written in pure Go, no CGO!



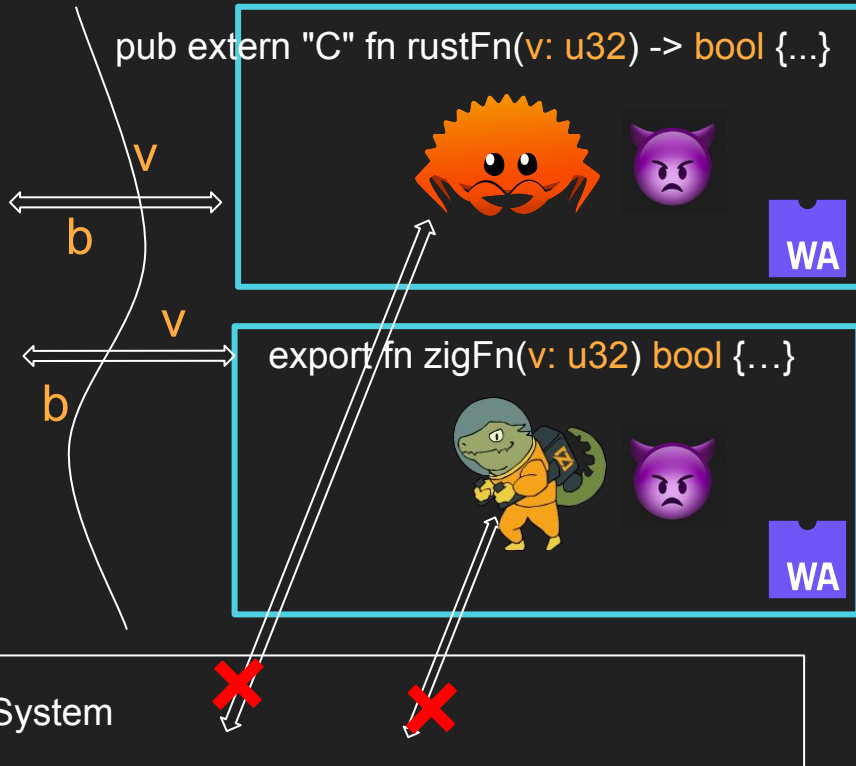
```
func main () {  
  b := rustFn(v) ↔  
  b = zigFn(v) ↔  
  ....  
}
```



?




```
func main () {  
  b := rustFn(v)  
  b = zigFn(v)  
  ....  
}
```





wazero

=> CGO-less Foreign Function Interface



FFI with wazero vs CGO

- No CGO
 - Static binary, cross compilation, etc
- Zero dependency
 - E.g. third party toolchains
- Compile once, run everywhere
- Sandbox environment
 - Memory isolation
 - Deny “system calls” by default



How it works: memory isolation



// Instantiate a Zig Wasm binary.

zig := r.InstantiateModuleFromBinary(...)



WA

Linear memory

// Instantiate a Rust Wasm binary.

rust := r.InstantiateModuleFromBinary(...)



WA

Linear memory




Go program



// Instantiate a Zig Wasm binary.

zig := r.InstantiateModuleFromBinary(...)




make([]byte, N)



Go program

// Instantiate a Rust Wasm binary.

rust := r.InstantiateModuleFromBinary(...)




make([]byte, M)



// Instantiate a Zig Wasm binary.

zig := r.InstantiateModuleFromBinary(...)




make([]byte, N)



Go program

// Instantiate a Rust Wasm binary.

rust := r.InstantiateModuleFromBinary(...)




make([]byte, M)



// Instantiate a Zig Wasm binary.


zig := r.InstantiateModuleFromBinary(...)



Go program

// Instantiate a Rust Wasm binary.

rust := r.InstantiateModuleFromBinary(...)



How it works: system call isolation



// Instantiate a Zig Wasm binary.

zig := r.InstantiateModuleFromBinary(...)



// Instantiate a Rust Wasm binary.

rust := r.InstantiateModuleFromBinary(...)



Linear memory

WA

Linear memory

WA

 `func readFile(fd int, ...) {...}`

`func writeFile(fd int, ...) {...}`

Go program

read(2)

foo.txt

Operating System

bar.txt

write(2)






// Instantiate a Zig Wasm binary.

zig := r.InstantiateModuleFromBinary(...)



// Instantiate a Rust Wasm binary.

rust := r.InstantiateModuleFromBinary(...)



func readFile(fd int, ...) {...}

Go program

func writeFile(fd int, ...) {...}

read(2)



Operating System



write(2)



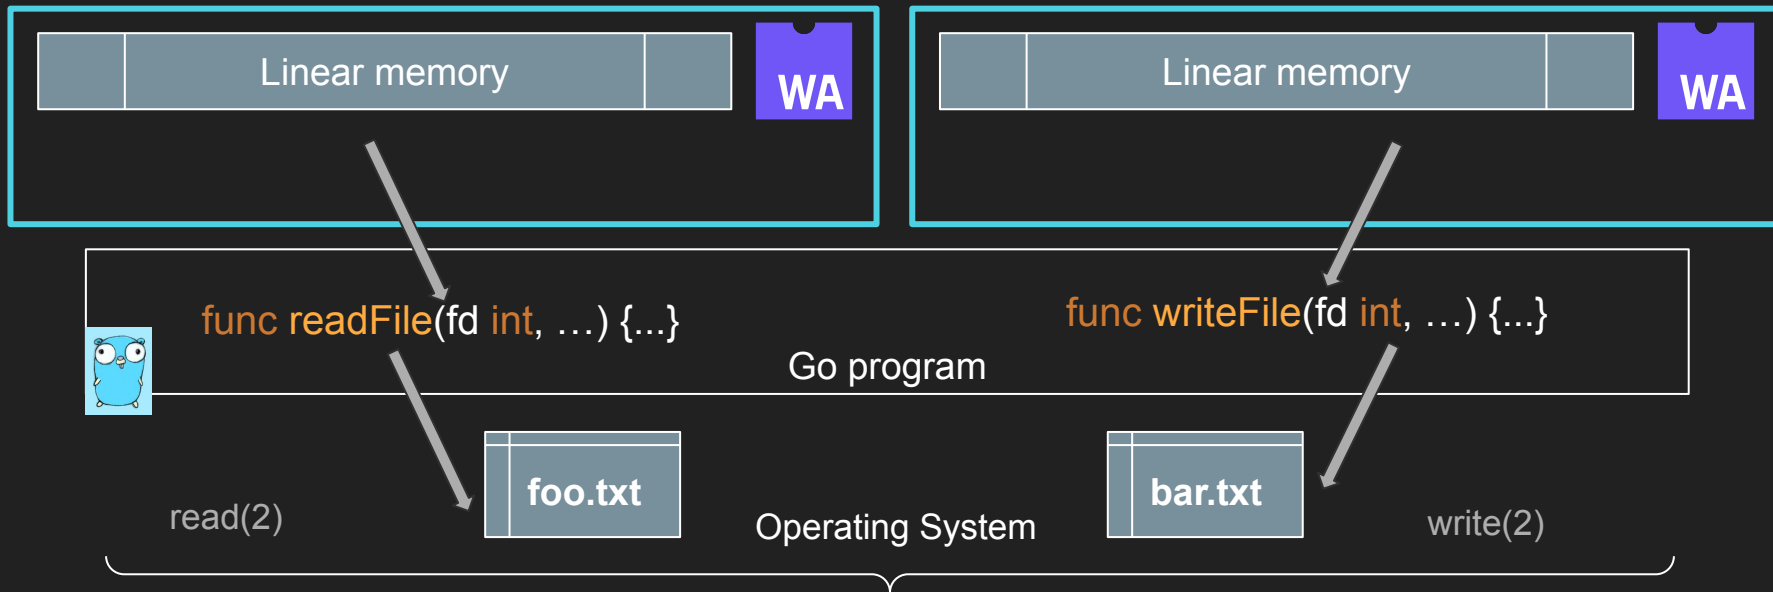
// Instantiate a Zig Wasm binary.

zig := r.InstantiateModuleFromBinary(...)



// Instantiate a Rust Wasm binary.

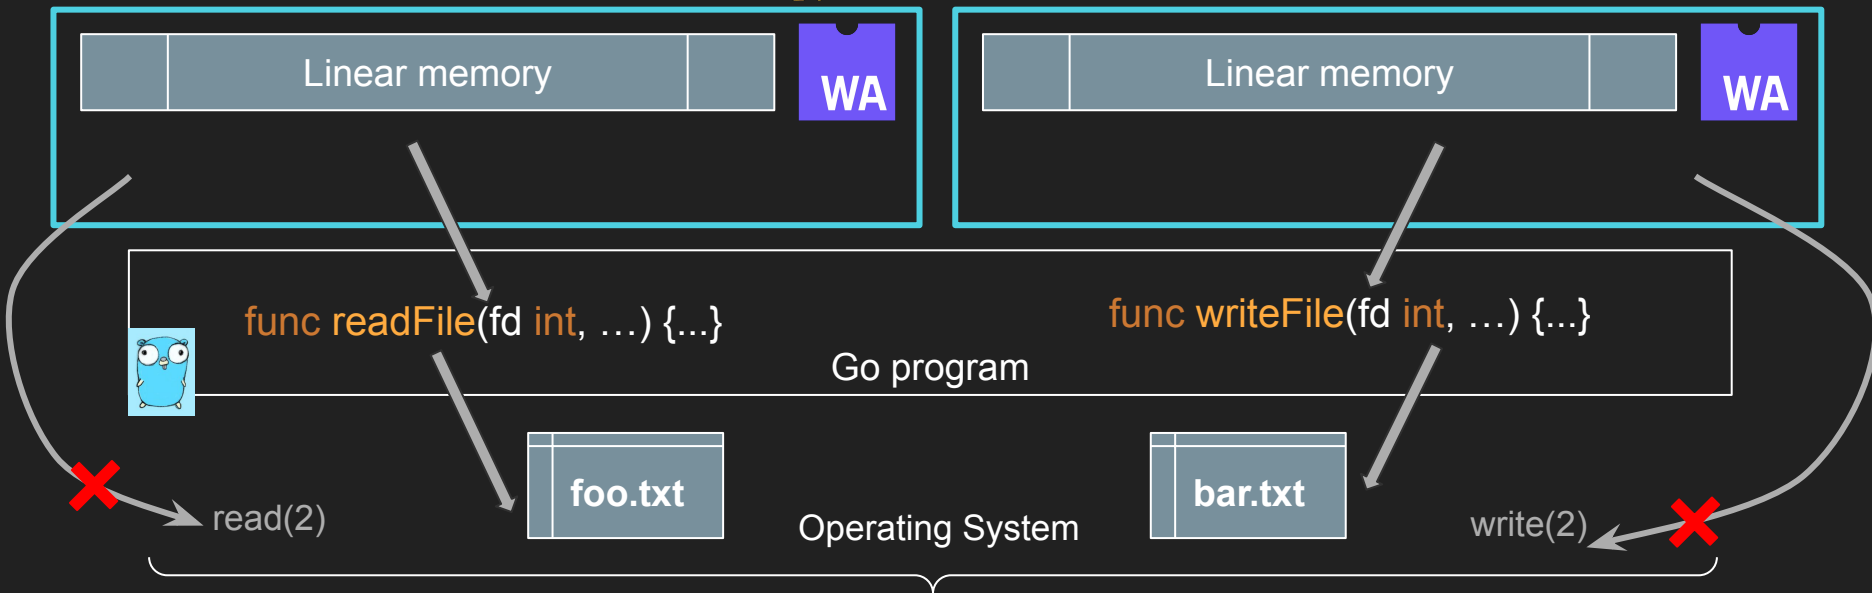
rust := r.InstantiateModuleFromBinary(...)



// Instantiate a Zig Wasm binary.
zig := r.InstantiateModuleFromBinary(...)



// Instantiate a Rust Wasm binary.
rust := r.InstantiateModuleFromBinary(...)



System Calls = Go functions
Memory = []byte{...}



// Create a new WebAssembly Runtime.

r := wazero.NewRuntime(ctx)

// Instantiate a Rust Wasm binary.

rust, _ := r.InstantiateModuleFromBinary(ctx, rustBinary) 

// Instantiate a Zig Wasm binary.

zig, _ := r.InstantiateModuleFromBinary(ctx, zigBinary)



// Call functions exported by Wasm modules.

... := rust.ExportedFunction("rustFn").Call(ctx, ...)

... := zig.ExportedFunction("zigFn").Call(ctx, ...)



Example projects!



- Trivy: vulnerability scanner
- Can extend scanning logics with Wasm, powered by wazero





dapr / dapr

Public



Watch 429



Fork 1.5k



Starred 19.3k



- dapr: portable, serverless application platform
- Can add HTTP middleware in Wasm powered by wazero





mathetake / wazero-sqlite

Public

Pin

Unwatch 3

Fork 0

Starred 5

- Running a Wasm-compiled SQLite inside Go, without CGO
- Possible implementation of CGO-less and sandboxed SQL Driver.





mathetake / wazero-re2

Public



Pin



Unwatch 1



Fork 0



Star 0

- re2: a fast regular expression engine in C++
- Running Wasm-compiled re2, without CGO
- In some cases, faster than regexp package in the Go std library!



Cons of FFI with wazero vs CGO

- Performance degradation
 - Wasm == Virtualization
 - Depends on runtime implementation
- Needs to compile your FFI to Wasm
 - Premature ecosystem
 - Refactor in a Wasm-friendly way



Cons of FFI with wazero vs CGO

- Performance degradation
 - Wasm == Virtualization
 - Depends on runtime implementation
- Needs to compile your FFI to Wasm
 - Premature ecosystem
 - Refactor in a Wasm-friendly way



Cons of FFI with wazero vs CGO

- Performance degradation
 - Wasm == Virtualization
 - Depends on runtime implementation
- Needs to compile your FFI to Wasm
 - Premature ecosystem
 - Refactor in a Wasm-friendly way



⚠ This commit does not belong to any branch on this repository, and may belong to a fork outside of the repository.

Add ability to build targeting wasi

Browse files



anuraaga committed 4 days ago

1 parent d61aa2e

commit 78f07ebbf92c164fcb9a5f7e13d0954a6eb01b47

Showing 6 changed files with 1,143 additions and 35 deletions.

Split

Unified

wazero deep dive...



Q. How correct is the implementation?



Q. How correct is the implementation?

A. 100% compatible with Wasm spec (1.0&2.0)



Q. How is wazero tested?



Q. How is wazero tested?

A. Specification tests & random binary fuzzing



Q. How is the VM implemented?



Q. How is the VM implemented?

A. Two modes: interpreter and AOT compiler



Interpreter mode

- Runs on any platform (GOOS/GOARCH)
- Fast startup time
- Slow execution

```
wazero.NewRuntimeConfigInterpreter()
```



Ahead-Of-Time (AOT) compiler mode

- Runs on {amd64,arm64} x {linux,darwin,windows,freebsd,etc}
- Slow startup time
 - AOT = compile Wasm binary into native machine code before execution
- Fast execution (10x+ faster than interpreter)

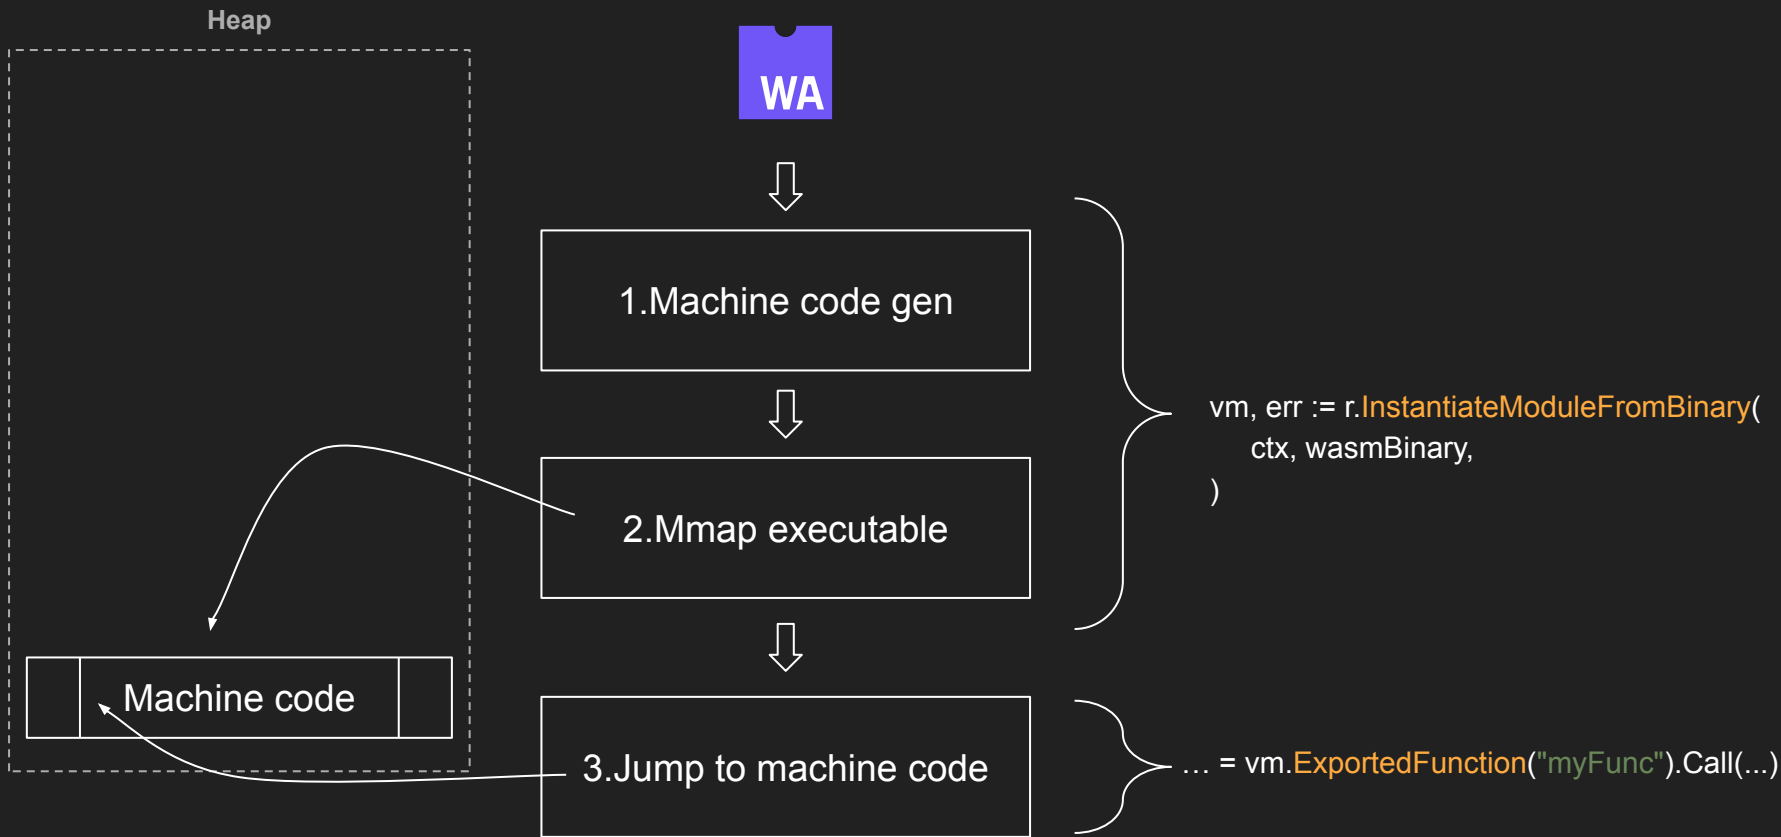
`wazero.NewRuntimeConfigCompiler()`



How AOT compiler works

1. Creates native machine code semantically equivalent to Wasm binary
2. mmap the machine code []byte as executable
3. Jumps into the “executable” []byte via a Go Assembly function





1. Machine code generation

```
VUMLAL: {u: 0b1, opcode: 0b1000, qAndSize: map[VectorArrangement]qAndSize{
    VectorArrangement2S: {q: 0b0, size: 0b10},
    VectorArrangement4H: {q: 0b0, size: 0b01},
    VectorArrangement8B: {q: 0b0, size: 0b00},
}},
SMULL: {u: 0b0, opcode: 0b1100, qAndSize: map[VectorArrangement]qAndSize{
    VectorArrangement8B: {q: 0b0, size: 0b00},
    VectorArrangement4H: {q: 0b0, size: 0b01},
    VectorArrangement2S: {q: 0b0, size: 0b10},
}},
SMULL2: {u: 0b0, opcode: 0b1100, qAndSize: map[VectorArrangement]qAndSize{
    VectorArrangement16B: {q: 0b1, size: 0b00},
    VectorArrangement8H: {q: 0b1, size: 0b01},
    VectorArrangement4S: {q: 0b1, size: 0b10},
}},
UMULL: {u: 0b1, opcode: 0b1100, qAndSize: map[VectorArrangement]qAndSize{
    VectorArrangement8B: {q: 0b0, size: 0b00},
    VectorArrangement4H: {q: 0b0, size: 0b01},
    VectorArrangement2S: {q: 0b0, size: 0b10},
}},
UMULL2: {u: 0b1, opcode: 0b1100, qAndSize: map[VectorArrangement]qAndSize{
    VectorArrangement16B: {q: 0b1, size: 0b00},
    VectorArrangement8H: {q: 0b1, size: 0b01},
    VectorArrangement4S: {q: 0b1, size: 0b10},
}},
```

Assembler



```
func (c *arm64Compiler) compileV128ExtMul(o *wazeroir.OperationV128ExtMul) error {
    var inst asm.Instruction
    var arr arm64.VectorArrangement
    if o.UseLow {
        if o.Signed {
            inst = arm64.SMULL
        } else {
            inst = arm64.UMULL
        }

        switch o.OriginShape {
        case wazeroir.ShapeI8x16:
            arr = arm64.VectorArrangement8B
        case wazeroir.ShapeI16x8:
            arr = arm64.VectorArrangement4H
        case wazeroir.ShapeI32x4:
            arr = arm64.VectorArrangement2S
        }
    } else {
        if o.Signed {
            inst = arm64.SMULL2
        } else {
            inst = arm64.UMULL2
        }
        arr = defaultArrangementForShape(o.OriginShape)
    }
    return c.compileV128x2BinOp(inst, arr)
}
```

Code generation

2. Mmap machine code as executable

```
func mmapCodeSegmentAMD64(code io.Reader, size int) ([]byte, error) {
    mmapFunc, err := syscall.Mmap(
        -1,
        0,
        size,
        // The region must be RWX: RW for writing native codes, X for executing the region.
        syscall.PROT_READ|syscall.PROT_WRITE|syscall.PROT_EXEC,
        // Anonymous as this is not an actual file, but a memory,
        // Private as this is in-process memory region.
        syscall.MAP_ANON|syscall.MAP_PRIVATE,
    )
    if err != nil {
        return nil, err
    }

    w := &bufWriter{underlying: mmapFunc}
    _, err = io.CopyN(w, code, int64(size))
    return mmapFunc, err
}
```



3. Jump into machine code

```
func (ce *callEngine) execWasmFunction(ctx context.Context, callCtx
    codeAddr := ce.initialFn.codeInitialAddress
    modAddr := ce.initialFn.moduleInstanceAddress

entry:
{
    // Call into the native code.
    nativecall(codeAddr, uintptr(unsafe.Pointer(ce)), modAddr)
```

```
#include "funcdata.h"
#include "textflag.h"
```



```
TEXT ·nativecall(SB),NOSPLIT,NOFRAME,$0-24
    MOVQ ce+8(FP),R13                // Load the address of *callEngine.
    MOVQ moduleInstanceAddress+16(FP),R12 // Load the address of *wasm.ModuleInstance
    MOVQ codeSegment+0(FP),AX        // Load the address of native code.
    JMP AX                          // Jump to native code.
```



Challenges in AOT compiler implementation

- Do not modify Goroutine-stack! (e.g. “call” instruction)
- Do not access Goroutine-stack allocated variable from machine code
- Debugging is extremely difficult
- Single pass compiler: optimizations are TODOs



Challenges in AOT compiler implementation

- Do not modify Goroutine-stack! (e.g. “call” instruction)
- Do not access Goroutine-stack allocated variable from machine code
- Debugging is extremely difficult
- Single pass compiler: optimizations are TODOs






[]byte{...} allocated by Go runtime

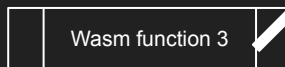
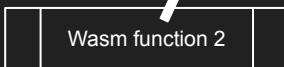
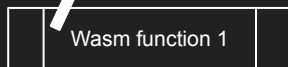
A struct allocated by Go runtime

Heap

[]byte{...} allocated by Go runtime



Function calls



....

Generated machine codes

Goroutine stack 1

....

Goroutine stack N



Challenges in AOT compiler implementation

- Do not modify Goroutine-stack! (e.g. “call” instruction)
- Do not access Goroutine-stack allocated variable from machine code
- Debugging is extremely difficult
- Naive single pass compiler: optimizations are TODOs



Wrap up!

- FFI == Calling non-Go functions from Go
- CGO works, but has some issues
- CGO-less FFI is possible with wazero+WebAssembly
- wazero is written in pure Go, zero dependency!



Thank you!

Twitter, GitHub: [@mathetake](#)

Gopher Slack: [#wazero](#)

